

# COP 701: Software Systems Lab

## (MCS192568 and MCS192574)

## Introduction

WebAssembly [1] (abbreviated Wasm) is a binary instruction format for a stack-based virtual machine. Wasm is designed as a portable target for compilation of high-level languages like C/C++/Rust, enabling deployment on the web for client and server applications.

## Deliverables - Part 1

### WebAssembly vs. Native Code

#### Experiment 1 - Relative Execution Time:

Run the PolyBench suite:

We have run all 30 programs of polyBench which are mentioned in benchmark\_list file in Polybench.



For native code we have used python scripts using **GCC** compiler

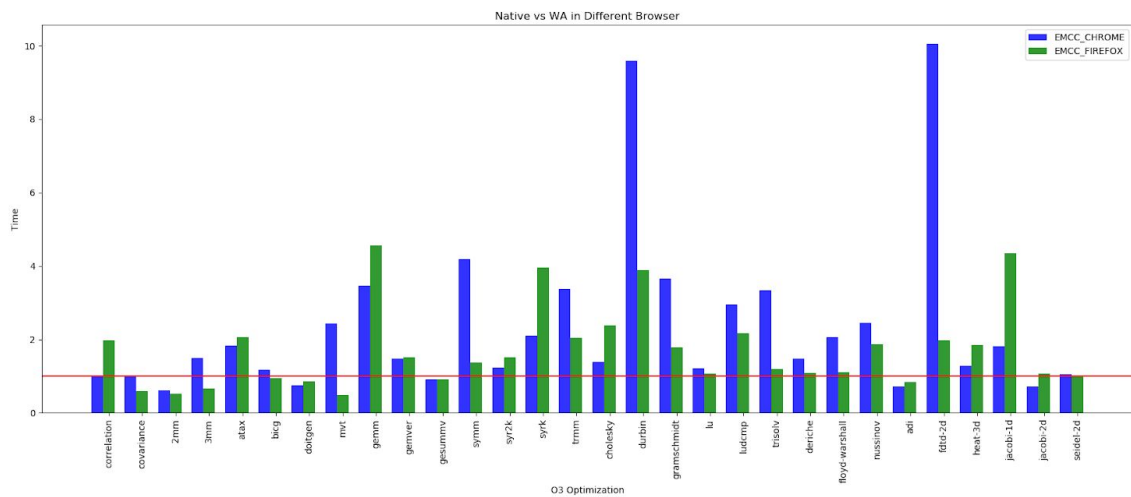
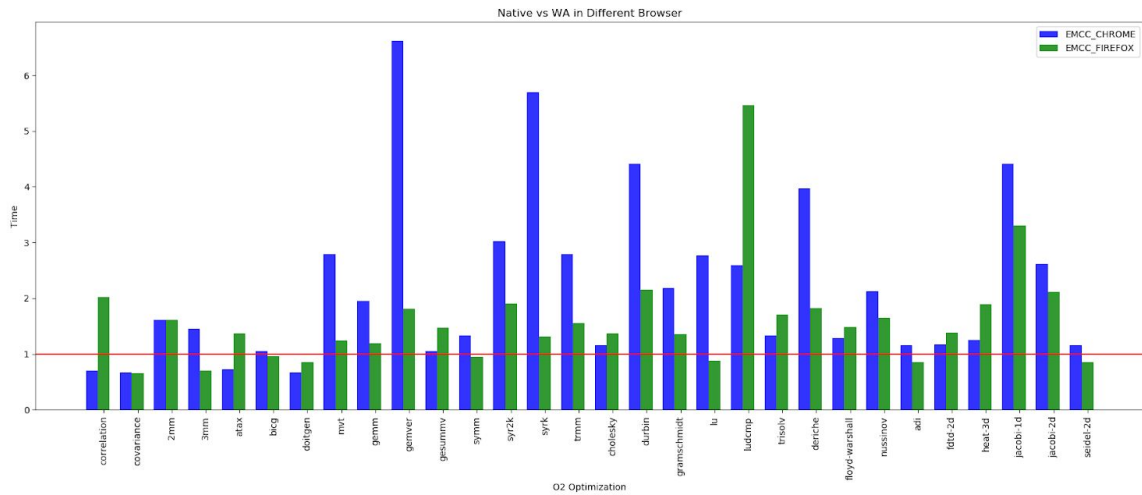
For Web assembly code we have used python scripts using **Emscripten** to generate the required html file to run on web.

We ran all files on **Chrome** and **Firefox** browser and noted the running time.

We stored the result in **PICKLE** format and used it in **plot1.py** to plot the graph using matplotlib. Pickles used :: **emcc.pickle** , **gcc.pickle**

Plot1.py is the python script to plot the graph using matplotlib





## **Conclusion:**

We can observe from the following graph the drastic improvement in GCC optimization from O0→ O3

It can be seen that as we optimize EMCC and GCC there is not much variance in run time of EMCC where as GCC shows a sharp drop in running time.

## TABLE DATA : NATIVE VS WA\_CHROME VS VS WA\_FIREFOX

| FileName       | Native   | Chrome | Firefox |
|----------------|----------|--------|---------|
| correlation    | 1.036447 | 0.385  | 0.238   |
| covariance     | 1.023256 | 0.353  | 0.24    |
| 2mm            | 0.217544 | 0.166  | 0.1     |
| 3mm            | 0.294644 | 0.191  | 0.131   |
| atax           | 0.067357 | 0.042  | 0.037   |
| bicg           | 0.094306 | 0.033  | 0.033   |
| dotgen         | 9.053019 | 2.822  | 2.975   |
| mvt            | 0.115363 | 0.076  | 0.074   |
| gemm           | 0.142328 | 0.087  | 0.034   |
| gemver         | 0.203771 | 0.102  | 0.097   |
| gesummv        | 0.039934 | 0.012  | 0.014   |
| symm           | 0.23491  | 0.052  | 0.032   |
| syr2k          | 0.45239  | 0.145  | 0.094   |
| syrk           | 0.128161 | 0.037  | 0.028   |
| trmm           | 0.120602 | 0.024  | 0.026   |
| cholesky       | 0.180052 | 0.057  | 0.05    |
| durbin         | 0.038307 | 0.024  | 0.031   |
| gramschmidt    | 0.204757 | 0.068  | 0.066   |
| lu             | 0.629148 | 0.164  | 0.112   |
| ludcmp         | 0.311057 | 0.14   | 0.131   |
| trisolv        | 0.018677 | 0.01   | 0.013   |
| deriche        | 1.700097 | 1.152  | 0.688   |
| floyd-warshall | 0.435825 | 0.165  | 0.088   |
| nussinov       | 2.19042  | 0.722  | 0.544   |
| adi            | 0.890928 | 0.233  | 0.175   |
| fttd-2d        | 0.421735 | 0.133  | 0.082   |
| heat-3d        | 1.00217  | 0.188  | 0.098   |
| jacobi-1d      | 0.023774 | 0.01   | 0.011   |
| jacobi-2d      | 0.24264  | 0.092  | 0.079   |
| seidel-2d      | 0.494345 | 0.322  | 0.306   |

| FileName       | Native   | Chrome | Firefox |
|----------------|----------|--------|---------|
| correlation    | 0.314698 | 0.284  | 0.321   |
| covariance     | 0.43711  | 0.269  | 0.175   |
| 2mm            | 0.101363 | 0.116  | 0.068   |
| 3mm            | 0.125874 | 0.177  | 0.12    |
| atax           | 0.031783 | 0.036  | 0.038   |
| bicg           | 0.02172  | 0.019  | 0.039   |
| dotgen         | 3.046613 | 2.687  | 2.848   |
| mvt            | 0.049865 | 0.049  | 0.044   |
| gemm           | 0.033279 | 0.059  | 0.088   |
| gemver         | 0.066835 | 0.135  | 0.091   |
| gesummv        | 0.010057 | 0.011  | 0.009   |
| symm           | 0.0298   | 0.044  | 0.034   |
| syr2k          | 0.060021 | 0.328  | 0.084   |
| syrk           | 0.026625 | 0.033  | 0.015   |
| trmm           | 0.026329 | 0.047  | 0.031   |
| cholesky       | 0.065887 | 0.12   | 0.055   |
| durbin         | 0.033482 | 0.018  | 0.032   |
| gramschmidt    | 0.08672  | 0.075  | 0.053   |
| lu             | 0.113472 | 0.22   | 0.127   |
| ludcmp         | 0.05554  | 0.067  | 0.108   |
| trisolv        | 0.010583 | 0.014  | 0.013   |
| deriche        | 0.691137 | 0.951  | 0.639   |
| floyd-warshall | 0.122582 | 0.135  | 0.12    |
| nussinov       | 0.325597 | 0.465  | 0.401   |
| adi            | 0.200365 | 0.402  | 0.181   |
| fttd-2d        | 0.074208 | 0.17   | 0.078   |
| heat-3d        | 0.09666  | 0.114  | 0.101   |
| jacobi-1d      | 0.002308 | 0.009  | 0.008   |
| jacobi-2d      | 0.033601 | 0.047  | 0.103   |
| seidel-2d      | 0.300143 | 0.267  | 0.256   |

| FileName       | Native   | Chrome | Firefox |
|----------------|----------|--------|---------|
| correlation    | 0.231351 | 0.163  | 0.467   |
| covariance     | 0.22522  | 0.151  | 0.148   |
| 2mm            | 0.045797 | 0.074  | 0.074   |
| 3mm            | 0.081769 | 0.119  | 0.058   |
| atax           | 0.037178 | 0.027  | 0.051   |
| bicg           | 0.020807 | 0.022  | 0.02    |
| doitgen        | 1.689575 | 1.128  | 1.44    |
| mvt            | 0.038663 | 0.108  | 0.048   |
| gemm           | 0.025122 | 0.049  | 0.03    |
| gemver         | 0.052417 | 0.347  | 0.095   |
| gesummv        | 0.009474 | 0.01   | 0.014   |
| symm           | 0.04421  | 0.059  | 0.042   |
| syr2k          | 0.066058 | 0.2    | 0.126   |
| syrk           | 0.015268 | 0.087  | 0.02    |
| trmm           | 0.022573 | 0.063  | 0.035   |
| cholesky       | 0.032887 | 0.038  | 0.045   |
| durbin         | 0.009279 | 0.041  | 0.02    |
| gramschmidt    | 0.042037 | 0.092  | 0.057   |
| lu             | 0.082995 | 0.23   | 0.073   |
| ludcmp         | 0.044794 | 0.116  | 0.245   |
| trisolv        | 0.005267 | 0.007  | 0.009   |
| deriche        | 0.633365 | 2.513  | 1.154   |
| floyd-warshall | 0.10943  | 0.141  | 0.163   |
| nussinov       | 0.308352 | 0.655  | 0.51    |
| adi            | 0.188833 | 0.218  | 0.162   |
| fttd-2d        | 0.063153 | 0.074  | 0.087   |
| heat-3d        | 0.091511 | 0.115  | 0.173   |
| jacobi-1d      | 0.004531 | 0.02   | 0.015   |
| jacobi-2d      | 0.032148 | 0.084  | 0.068   |
| seidel-2d      | 0.297752 | 0.347  | 0.256   |

| FileName       | Native   | Chrome | Firefox |
|----------------|----------|--------|---------|
| correlation    | 0.162554 | 0.161  | 0.32    |
| covariance     | 0.204879 | 0.201  | 0.121   |
| 2mm            | 0.079362 | 0.049  | 0.041   |
| 3mm            | 0.086416 | 0.129  | 0.058   |
| atax           | 0.016999 | 0.031  | 0.035   |
| bicg           | 0.021131 | 0.025  | 0.02    |
| doitgen        | 1.794551 | 1.344  | 1.542   |
| mvt            | 0.080292 | 0.196  | 0.039   |
| gemm           | 0.011829 | 0.041  | 0.054   |
| gemver         | 0.054802 | 0.081  | 0.083   |
| gesummv        | 0.009787 | 0.009  | 0.009   |
| symm           | 0.034428 | 0.144  | 0.047   |
| syr2k          | 0.10422  | 0.128  | 0.158   |
| syrk           | 0.016662 | 0.035  | 0.066   |
| trmm           | 0.013661 | 0.046  | 0.028   |
| cholesky       | 0.025218 | 0.035  | 0.06    |
| durbin         | 0.007197 | 0.069  | 0.028   |
| gramschmidt    | 0.038742 | 0.142  | 0.069   |
| lu             | 0.091777 | 0.112  | 0.099   |
| ludcmp         | 0.055867 | 0.165  | 0.121   |
| trisolv        | 0.00748  | 0.025  | 0.009   |
| deriche        | 0.586128 | 0.867  | 0.637   |
| floyd-warshall | 0.108747 | 0.225  | 0.12    |
| nussinov       | 0.289256 | 0.709  | 0.54    |
| adi            | 0.195417 | 0.139  | 0.163   |
| fttd-2d        | 0.033311 | 0.335  | 0.066   |
| heat-3d        | 0.081371 | 0.104  | 0.15    |
| jacobi-1d      | 0.002757 | 0.005  | 0.012   |
| jacobi-2d      | 0.061438 | 0.044  | 0.066   |
| seidel-2d      | 0.251114 | 0.263  | 0.255   |

## **Experiment 2 - Instruction Mix Analysis:**

### **Pin Tool:**

We used the pin tool ver-3.11 to benchmark the native code.

For Instruction breakup we modified the **inscount0.cpp**

**Pin2\_3.py** is the name of python script used to run the benchmark

The namespaces mentioned in PIN manual are used to count different breakups of instruction are present in inscount file.

The result observed is stored as PICKLE format to be used in **plot2.py** python script which generates bar graph.

### **Wasabi tool**



[5] to instrument the WebAssembly code to obtain the instruction breakup. Types: load, store, arithmetic/logical-int, arithmetic/logical- oat, branch, register transfer, And Nop



We use all the hooks to get total instruction count as well as different types of instructions.

We have made a thread script which parallely use the emcc for different browsers and gcc for native machine in question 1 and parallely spawns emrun thread for question 2.

We modified the js file **instruction-mix.js** to find total instruction count, breakup instruction count and frequency of function executions. To prevent manual intervention we have used a POST call to the emrun server to post the output on server console and use subprocess.Popen of python to get the data.

**Plot2.py** is the python script to plot the graph using matplotlib





## Conclusion:

Improvements can be observed in both NATIVE and WEBASSEMBLY calculated using PIN and WASABI tool for instrumentation.

In native machine it is observed that as we increase the optimization, more instructions are incorporated into the above mentioned type.



The convergence of type of instruction into our used format is comparatively slow in WEBASSEMBLY.

### Experiment 3 - Instruction Count:

For native code we again used PIN tool to count no of dynamic instruction count using the following namespace:


**Pin2\_3.py** is the name of python script used to run the benchmark


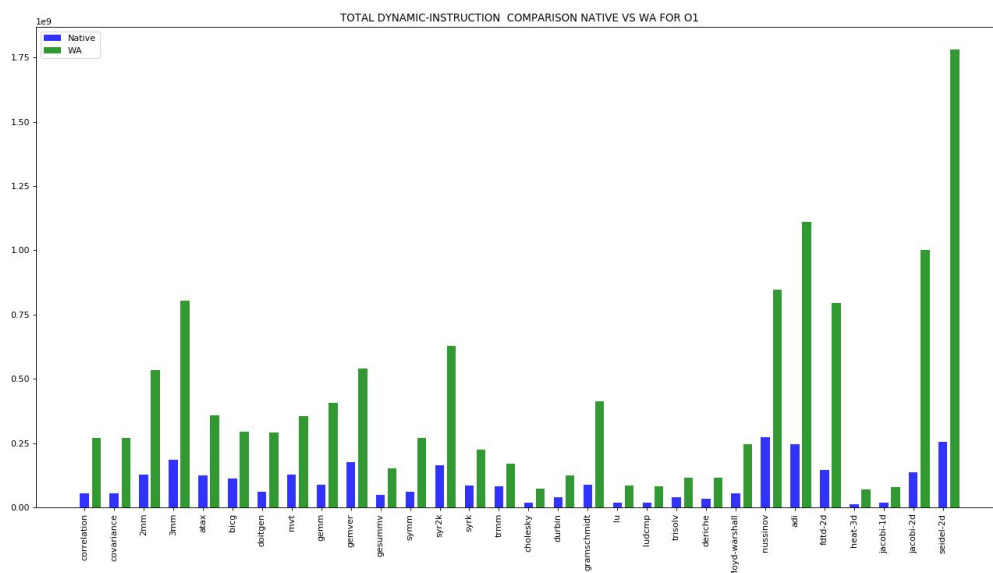
**Plot3.py** python script is used to plot the graph showing relative


Plot a bar graph showing the

relative number of dynamic instructions for WebAssembly vis-a-vis Native Code.

**Plot3.py** is the python script to plot the graph using matplotlib







## Conclusion

As we Optimise our GCC there is significant drop in total number of Dynamic instructions which supports the observations of experiment 1 that total time reduces in majority cases.

In output of WebAssembly instrumented by wasabi is not that impressive as native machine.

• **Experiment 4 - Hot Code:** Identify the functions in the code that are run frequently and plot their dynamic instruction count.

**PIN:** Tool is used to count the Hot code by sorting the functions in descending order according to number of calls per function.

The name and ratio of count of INstruction of each function is sent as output.


The result is stored in pickle format as **nativePin4.pickle**



A Vector< Pair <FunName,InsCount>, CallCount > is used to calculate the required instruction.


**WASABI:**

In the hot code section we have used stack implementation to keep track of all the dynamic instructions inside the function. Additionally we made sure that if a function A gets called inside function B then instruction count of function A should not include instruction count of function B. Additionally we used wasabi.wast file which is generated by using -g3 flag in emcc command to map function name to its function number.

The `call_pre` hook was used to detect the start of a function whereas `call_post` was used to detect the end of a function. So we pushed a function on stack at `call_pre` hook and popped at `call_post` hook and in between all the instructions were used to increment instruction count of the function on top of the stack.







## Conclusion

The above bar graph plots the relative hotcode in each function.

Since we count the total Instruction of top 4 Hotcodes and rest are included in others hence relative value of “others count” to “hot codes count” is very huge sometimes tending to 99% hence we have scaled the y axis in bar graph to log.

Since we sort the functions as per their calls rather than Instruction code, hence function with large instruction count is included in others. Making it the most dominant category even though it does not contain top hot codes.

# Native Hardware

|                 |   |
|-----------------|---|
| DRAM Size       | 12 GB DDR3  |
| CPU Frequency   | Intel(R) Core(TM) i3-5010U CPU @ 2.10GHz (i3 - 5th Generation)    |
| Cache Size      | 3072 KB   |
| CPU cores       | 2   |
| OS Version      | Ubuntu 16.04.6 LTS (codename : Xenial)<br>with GNOME Shell 3.18.5 |
| GCC Version     | gcc (Ubuntu 5.4.0-6ubuntu1~16.04.11)<br>5.4.0 20160609            |
| Python Version  | Python 3.7.3  |
| Google Chrome   | Google Chrome 77.0.3865.120                                       |
| Mozilla Firefox | Mozilla Firefox 69.0.1  |
| PolyBench       | polybench-c-4.1   |